

Small-Data Surrogate Shock-Aware Designed (S3) models: Gradient-Free Residual Adaptation with Adaptive Conformal Uncertainty for Engineering Time Series

A.patino

Abstract

F

Keywords: small-data forecasting, residual adaptation, reservoir computing, conformal prediction, non-stationarity, industrial forecasting

1. Introduction

1.1. Key Contributions

1. We formulate small-data engineering forecasting as a *surrogate residual adaptation* problem and propose a gradient-free correction architecture that freezes the baseline forecaster and learns only a low-capacity residual readout, thereby prioritizing variance control and convex optimization stability over end-to-end parameter proliferation [7, 13].
2. We introduce two complementary model realizations—*S3-Forecaster*, based on fixed-reservoir Echo State dynamics under spectral radius constraints, and *S3-FastSketch*, based on deterministic causal multiscale sketches—to separate temporal representation from lightweight residual decoding in genuinely short-series regimes [41, 42, 44].
3. We develop a shock-aware residual control mechanism based on out-of-bag gating and bounded shrinkage, motivated by the documented adaptation lag of classical online updaters under concept drift and by the need to suppress variance amplification when apparent residual structure is not stable enough to be learnable [25, 26].
4. We integrate a sequential Adaptive Conformal Inference layer for uncertainty quantification, so that prediction intervals are calibrated by recent nonconformity scores and retain distribution-free predictive validity despite exchangeability violation in dependent engineering sequences [29, 36, 37].
5. We adopt a strict chronological train–calibration–test protocol designed to preserve temporal order and avoid information leakage and look-ahead bias, aligning evaluation with the actual decision task faced in forecast deployment rather than with i.i.d. cross-validation surrogates [39, 40].

2. Materials and Methods

2.1. Problem Formulation

Let

$$\mathbf{y}_{1:N} = (y_1, \dots, y_N)^\top \in \mathbb{R}^N \quad (1)$$

denote a univariate engineering time series with monthly sampling frequency and limited historical availability, $N < 200$. The forecasting problem is defined with respect to the natural filtration

$$\mathcal{F}_t = \sigma(y_1, \dots, y_t), \quad (2)$$

which represents all target information observed up to time t . A valid one-step-ahead predictor must satisfy

$$\hat{y}_{t|t-1} \in \mathcal{F}_{t-1}, \quad (3)$$

such that y_t cannot influence its own forecast. The proposed family follows a residual-adaptation decomposition:

$$\hat{y}_{t|t-1} = \hat{y}_{t|t-1}^{\text{base}} + \hat{\Delta}_{t|t-1}, \quad (4)$$

where $\hat{y}_{t|t-1}^{\text{base}}$ is generated by a frozen causal prior and $\hat{\Delta}_{t|t-1}$ is a data-efficient residual correction. The causal base residual is

$$e_t = y_t - \hat{y}_{t|t-1}^{\text{base}}. \quad (5)$$

The proposed models differ only in the representation used to estimate the predictable component of e_t :

$$\text{S3-Forecaster: } e_{1:t-1} \mapsto \mathbf{h}_{t-1} \mapsto \hat{r}_{t|t-1}, \quad (6)$$

$$\text{S3-FastSketch: } e_{1:t-1} \mapsto \phi_{t-1} \mapsto \hat{r}_{t|t-1}. \quad (7)$$

Neither architecture backpropagates through the causal prior or through its nonlinear feature extractor. The learned component is restricted to a regularized linear readout and a low-dimensional calibration mechanism.

2.2. Frozen Causal Forecast Prior

Let

$$f_{\text{base}} : \bigcup_{t \geq 1} \mathbb{R}^t \longrightarrow \mathbb{R} \quad (8)$$

denote a statistical, foundation-model, or heuristic forecasting operator whose parameters are frozen during residual adaptation. Its strictly causal output is

$$\hat{y}_{t|t-1}^{\text{base}} = f_{\text{base}}(y_1, \dots, y_{t-1}; \boldsymbol{\theta}_{\text{base}}). \quad (9)$$

The generic formulation permits the use of rolling estimators, ETS, Prophet, ARIMA, or a zero-shot time-series foundation model. The residual adapter does not alter $\boldsymbol{\theta}_{\text{base}}$. Consequently, adaptation cost is

independent of the number of parameters in the frozen prior. For the lightweight default prior, a causal rolling level estimate is defined as

$$\hat{\mu}_{t-1}^{(w)} = \frac{1}{|\mathcal{W}_{t-1}|} \sum_{j \in \mathcal{W}_{t-1}} y_j, \quad \mathcal{W}_{t-1} = \{\max(1, t-w), \dots, t-1\}. \quad (10)$$

When sufficient monthly history is available, a causal seasonal correction may be added:

$$\hat{s}_{m(t)} = \frac{1}{|\mathcal{I}_{m(t)}|} \sum_{j \in \mathcal{I}_{m(t)}} \left(y_j - \hat{\mu}_{j-1}^{(w)} \right), \quad (11)$$

where $\mathcal{I}_{m(t)}$ contains only pre- t observations belonging to the same calendar month as t . The resulting prior is

$$\hat{y}_{t|t-1}^{\text{base}} = \hat{\mu}_{t-1}^{(w)} + \hat{s}_{m(t)}. \quad (12)$$

2.3. S3-Forecaster: Fixed Reservoir Residual Adaptation

2.3.1. Residual Normalization

The residual stream is standardized using statistics estimated exclusively from the training block:

$$\tilde{e}_t = \frac{e_t - \hat{m}_e}{\hat{q}_{0.75}(e) - \hat{q}_{0.25}(e) + \varepsilon_s}, \quad t \in \mathcal{D}_{\text{tr}}, \quad (13)$$

where $\hat{m}_e$ is the training median and $\varepsilon_s > 0$ prevents division by zero. The same frozen transformation is used for calibration and testing.

2.3.2. Echo State Representation

Let R denote the reservoir dimension. The recurrent state is

$$\mathbf{h}_t \in [-1, 1]^R. \quad (14)$$

A leaky ESN update is defined by

$$\tilde{\mathbf{h}}_t = \tanh \left(\mathbf{W}_{\text{in}} \begin{bmatrix} 1 \\ \tilde{e}_t \end{bmatrix} + \mathbf{W}_{\text{res}} \mathbf{h}_{t-1} \right), \quad (15)$$

$$\mathbf{h}_t = (1 - \lambda_{\text{leak}}) \mathbf{h}_{t-1} + \lambda_{\text{leak}} \tilde{\mathbf{h}}_t, \quad \lambda_{\text{leak}} \in (0, 1]. \quad (16)$$

The fixed matrices satisfy

$$\mathbf{W}_{\text{in}} \in \mathbb{R}^{R \times 2}, \quad (17)$$

$$\mathbf{W}_{\text{res}} \in \mathbb{R}^{R \times R}. \quad (18)$$

They are sampled once from seed-controlled random distributions and are not updated during training. Let $\widetilde{\mathbf{W}}_{\text{res}}$ be an initially sampled recurrent matrix. Spectral normalization is performed as

$$\mathbf{W}_{\text{res}} = \rho_{\star} \frac{\widetilde{\mathbf{W}}_{\text{res}}}{\rho(\widetilde{\mathbf{W}}_{\text{res}}) + \varepsilon_{\rho}}, \quad 0 < \rho_{\star} < 1.5, \quad (19)$$

where $\rho(\cdot)$ denotes spectral radius. The stronger contraction condition

$$\|\mathbf{W}_{\text{res}}\|_2 < 1.5 \quad (20)$$

is sufficient for asymptotic forgetting of initial states under a 1-Lipschitz activation. In practice, $\rho(\mathbf{W}_{\text{res}}) < 1.5$ is used as the standard reservoir-design heuristic. Its purpose is to prevent uncontrolled recurrent amplification while retaining a fading memory of the residual history. To supplement the nonlinear reservoir with explicit short-term information, an autoregressive residual vector is constructed:

$$\boldsymbol{\ell}_t^{(p)} = [e_t, e_{t-1}, \dots, e_{t-p+1}]^\top \in \mathbb{R}^p. \quad (21)$$

The complete S3 representation is

$$\mathbf{z}_t = \left[1, \mathbf{h}_t^\top, \left(\boldsymbol{\ell}_t^{(p)} \right)^\top \right]^\top \in \mathbb{R}^P, \quad P = 1 + R + p. \quad (22)$$

2.3.3. Bayesian Ridge Residual Readout

For one-step-ahead forecasting, the readout assumes

$$e_{t+1} = \mathbf{w}^\top \mathbf{z}_t + \xi_{t+1}, \quad \xi_{t+1} \sim \mathcal{N}(0, \beta^{-1}). \quad (23)$$

A zero-mean isotropic Gaussian prior is imposed:

$$p(\mathbf{w} \mid \alpha) = \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I}_P). \quad (24)$$

Let

$$\mathbf{Z} = [\mathbf{z}_{t_1}, \dots, \mathbf{z}_{t_n}]^\top \in \mathbb{R}^{n \times P}, \quad (25)$$

$$\mathbf{r} = [e_{t_1+1}, \dots, e_{t_n+1}]^\top \in \mathbb{R}^n. \quad (26)$$

The negative log-posterior, excluding constants independent of $\mathbf{w}$, is

$$\mathcal{J}(\mathbf{w}) = \frac{\beta}{2} \|\mathbf{r} - \mathbf{Z}\mathbf{w}\|_2^2 + \frac{\alpha}{2} \|\mathbf{w}\|_2^2. \quad (27)$$

The Hessian is

$$\nabla_{\mathbf{w}}^2 \mathcal{J}(\mathbf{w}) = \beta \mathbf{Z}^\top \mathbf{Z} + \alpha \mathbf{I}_P \succ 0 \quad (28)$$

for $\alpha > 0$. Therefore, the objective is strictly convex and has a unique global minimizer. Setting the gradient to zero gives

$$\nabla_{\mathbf{w}} \mathcal{J}(\mathbf{w}) = -\beta \mathbf{Z}^\top (\mathbf{r} - \mathbf{Z}\mathbf{w}) + \alpha \mathbf{w} = \mathbf{0}, \quad (29)$$

$$(\beta \mathbf{Z}^\top \mathbf{Z} + \alpha \mathbf{I}_P) \mathbf{w} = \beta \mathbf{Z}^\top \mathbf{r}. \quad (30)$$

Hence,

$$\hat{\mathbf{w}}_{\text{MAP}} = (\beta \mathbf{Z}^\top \mathbf{Z} + \alpha \mathbf{I}_P)^{-1} \beta \mathbf{Z}^\top \mathbf{r}. \quad (31)$$

Equivalently, defining $\lambda_R = \alpha/\beta$,

$$\hat{\mathbf{w}}_{\text{MAP}} = (\mathbf{Z}^\top \mathbf{Z} + \lambda_R \mathbf{I}_P)^{-1} \mathbf{Z}^\top \mathbf{r}. \quad (32)$$

The complete posterior is

$$p(\mathbf{w} \mid \mathbf{r}, \mathbf{Z}, \alpha, \beta) = \mathcal{N}(\boldsymbol{\mu}_w, \boldsymbol{\Sigma}_w), \quad (33)$$

$$\boldsymbol{\Sigma}_w = (\alpha \mathbf{I}_P + \beta \mathbf{Z}^\top \mathbf{Z})^{-1}, \quad (34)$$

$$\boldsymbol{\mu}_w = \beta \boldsymbol{\Sigma}_w \mathbf{Z}^\top \mathbf{r}. \quad (35)$$

The one-step residual correction is

$$\hat{r}_{t+1|t} = \boldsymbol{\mu}_w^\top \mathbf{z}_t. \quad (36)$$

The Bayesian derivation provides a probabilistic interpretation of the regularized readout. When fixed precision parameters are used, its posterior mode is numerically equivalent to conventional Ridge regression. Lasso and Elastic Net may be considered as convex sensitivity variants, but they must not be described as Bayesian Ridge unless the corresponding prior interpretation is explicitly introduced.

2.3.4. Volatility-Sensitive OOB Gate

The residual correction is not assumed to be uniformly reliable. A causal volatility process is therefore computed from the prior residuals. Let

$$\mu_t^{(e)} = (1 - \lambda_v) \mu_{t-1}^{(e)} + \lambda_v e_t, \quad (37)$$

$$q_t^{(e)} = (1 - \lambda_v) q_{t-1}^{(e)} + \lambda_v \left(e_t - \mu_{t-1}^{(e)} \right)^2, \quad (38)$$

$$v_t = \sqrt{q_t^{(e)} + \varepsilon_v}. \quad (39)$$

The volatility is robustly standardized using training-only statistics:

$$\tilde{v}_t = \frac{v_t - \text{median}_{j \in \mathcal{D}_{\text{tr}}}(v_j)}{\text{IQR}_{j \in \mathcal{D}_{\text{tr}}}(v_j)/1.349 + \varepsilon_v}. \quad (40)$$

The gate is

$$g_t(a, b) = \sigma(a(\tilde{v}_t - b)) = \frac{1}{1 + \exp[-a(\tilde{v}_t - b)]}, \quad g_t \in (0, 1). \quad (41)$$

The ESN and its readout are fitted on $\mathcal{D}_{\text{tr}}$. Keeping these components fixed, the two gate parameters are estimated exclusively on the chronological calibration block:

$$(\hat{a}, \hat{b}) = \arg \min_{a \in [a_{\min}, a_{\max}], b \in [b_{\min}, b_{\max}]} \frac{1}{|\mathcal{D}_{\text{cal}}|} \sum_{t \in \mathcal{D}_{\text{cal}}} |e_t - g_{t-1}(a, b) \hat{r}_{t|t-1}|. \quad (42)$$

This is an out-of-bag optimization with respect to the readout because no calibration observation is used to estimate $\mathbf{w}$. The gate is a two-parameter calibration mechanism rather than a second high-capacity predictor. The final S3 point forecast is

$$\hat{y}_{t|t-1}^{\text{S3}} = \hat{y}_{t|t-1}^{\text{base}} + g_{t-1}(\hat{a}, \hat{b}) \hat{r}_{t|t-1}. \quad (43)$$

The term *gradient-free* refers to the absence of end-to-end backpropagation through the causal prior and reservoir. The bounded low-dimensional calibration of (a, b) is not conflated with neural gradient-based fine-tuning.

2.4. S3-FastSketch: Lightweight Multiscale Residual Adaptation

2.4.1. Causal Multiscale Feature Map

S3-FastSketch replaces the recurrent reservoir by a deterministic causal map

$$\Phi : e_{1:t} \mapsto \phi_t \in \mathbb{R}^{D_\phi}. \quad (44)$$

The representation is

$$\phi_t = [\ell_t^\top, \mathbf{a}_t^\top, \mathbf{v}_t^\top, \mathbf{d}_t^\top, \mathbf{c}_t^\top, \mathbf{s}_t^\top]^\top. \quad (45)$$

Let $\mathcal{Q} = \{q_1, \dots, q_Q\}$ be a set of causal smoothing spans and $\mathcal{K} = \{k_1, \dots, k_K\}$ a set of local filter scales.

Residual lag block.. For p autoregressive lags,

$$\ell_t = [e_{t-1}, e_{t-2}, \dots, e_{t-p}]^\top \in \mathbb{R}^p. \quad (46)$$

Exponentially weighted residual block.. For each $q \in \mathcal{Q}$,

$$a_t^{(q)} = (1 - \eta_q) a_{t-1}^{(q)} + \eta_q e_t, \quad (47)$$

$$\eta_q = \frac{2}{q+1}. \quad (48)$$

The block is

$$\mathbf{a}_t = [a_t^{(q_1)}, \dots, a_t^{(q_Q)}]^\top. \quad (49)$$

Residual-magnitude block.. The multiscale magnitude summaries are

$$v_t^{(q)} = (1 - \eta_q) v_{t-1}^{(q)} + \eta_q |e_t|, \quad (50)$$

and

$$\mathbf{v}_t = [v_t^{(q_1)}, \dots, v_t^{(q_Q)}]^\top. \quad (51)$$

Local-difference block.. For $q \in \mathcal{Q}$,

$$d_t^{(q)} = \begin{cases} e_t - e_{t-q}, & t > q, \\ 0, & t \leq q. \end{cases} \quad (52)$$

Thus,

$$\mathbf{d}_t = [d_t^{(q_1)}, \dots, d_t^{(q_Q)}]^\top. \quad (53)$$

Fixed causal convolutional block.. For each $k \in \mathcal{K}$, define the causal window

$$\mathcal{W}_{t,k} = \{\max(1, t - k + 1), \dots, t\}. \quad (54)$$

Its local mean is

$$c_{t,k}^{(\text{mean})} = \frac{1}{|\mathcal{W}_{t,k}|} \sum_{j \in \mathcal{W}_{t,k}} e_j. \quad (55)$$

Its endpoint slope is

$$c_{t,k}^{(\text{slope})} = e_t - e_{\max(1, t-k+1)}. \quad (56)$$

Let $\mathcal{W}_{t,k}^L$ and $\mathcal{W}_{t,k}^R$ denote the first and second halves of $\mathcal{W}_{t,k}$. The local contrast is

$$c_{t,k}^{(\text{contrast})} = \frac{1}{|\mathcal{W}_{t,k}^R|} \sum_{j \in \mathcal{W}_{t,k}^R} e_j - \frac{1}{|\mathcal{W}_{t,k}^L|} \sum_{j \in \mathcal{W}_{t,k}^L} e_j. \quad (57)$$

The local absolute energy is

$$c_{t,k}^{(\text{energy})} = \frac{1}{|\mathcal{W}_{t,k}|} \sum_{j \in \mathcal{W}_{t,k}} |e_j|. \quad (58)$$

The convolutional block is the concatenation

$$\mathbf{c}_t = \text{concat}_{k \in \mathcal{K}} \left[c_{t,k}^{(\text{mean})}, c_{t,k}^{(\text{slope})}, c_{t,k}^{(\text{contrast})}, c_{t,k}^{(\text{energy})} \right]. \quad (59)$$

Seasonal and calendar block.. For monthly seasonality $m = 12$,

$$\mathbf{s}_t = \left[e_{t-m}, e_t - e_{t-m}, \sin\left(\frac{2\pi t}{m}\right), \cos\left(\frac{2\pi t}{m}\right) \right]^\top, \quad (60)$$

where unavailable seasonal lags are initialized to zero. Each feature dimension is standardized using training medians and interquartile ranges:

$$\tilde{\phi}_{t,j} = \frac{\phi_{t,j} - \text{median}_{u \in \mathcal{D}_{\text{tr}}}(\phi_{u,j})}{\text{IQR}_{u \in \mathcal{D}_{\text{tr}}}(\phi_{u,j})/1.349 + \varepsilon_\phi}. \quad (61)$$

2.4.2. Regularized Residual Readout

For a forecast horizon H , define

$$\mathbf{r}_{t+1:t+H} = [e_{t+1}, \dots, e_{t+H}]^\top \in \mathbb{R}^H. \quad (62)$$

The multivariate residual readout is

$$\hat{\mathbf{r}}_{t+1:t+H|t} = \mathbf{W}^\top \tilde{\boldsymbol{\phi}}_t, \quad \mathbf{W} \in \mathbb{R}^{D_\phi \times H}. \quad (63)$$

Let

$$\boldsymbol{\Phi} \in \mathbb{R}^{n_{\text{tr}} \times D_\phi}, \quad (64)$$

$$\mathbf{R} \in \mathbb{R}^{n_{\text{tr}} \times H} \quad (65)$$

contain the training sketches and future residual targets. The readout solves

$$\widehat{\mathbf{W}} = \arg \min_{\mathbf{W}} \|\mathbf{R} - \boldsymbol{\Phi} \mathbf{W}\|_F^2 + \lambda_F \|\mathbf{W}\|_F^2. \quad (66)$$

Its closed-form solution is

$$\widehat{\mathbf{W}} = (\boldsymbol{\Phi}^\top \boldsymbol{\Phi} + \lambda_F \mathbf{I}_{D_\phi})^{-1} \boldsymbol{\Phi}^\top \mathbf{R}. \quad (67)$$

2.4.3. Projected Calibration Shrinkage

The raw residual forecast may be unstable when the calibration block is small. For each horizon h , a single bounded coefficient is therefore estimated from calibration pairs

$$\{(r_{i,h}, \hat{r}_{i,h})\}_{i \in \mathcal{D}_{\text{cal}}} . \quad (68)$$

The calibration objective is

$$\mathcal{L}_h(\gamma) = \sum_{i \in \mathcal{D}_{\text{cal}}} (r_{i,h} - \gamma \hat{r}_{i,h})^2 . \quad (69)$$

Differentiating gives

$$\frac{\partial \mathcal{L}_h}{\partial \gamma} = -2 \sum_{i \in \mathcal{D}_{\text{cal}}} \hat{r}_{i,h} (r_{i,h} - \gamma \hat{r}_{i,h}) , \quad (70)$$

$$0 = \sum_{i \in \mathcal{D}_{\text{cal}}} \hat{r}_{i,h} r_{i,h} - \gamma \sum_{i \in \mathcal{D}_{\text{cal}}} \hat{r}_{i,h}^2 . \quad (71)$$

The unconstrained solution is

$$\tilde{\gamma}_h = \frac{\sum_{i \in \mathcal{D}_{\text{cal}}} \hat{r}_{i,h} r_{i,h}}{\sum_{i \in \mathcal{D}_{\text{cal}}} \hat{r}_{i,h}^2 + \varepsilon_\gamma} . \quad (72)$$

The final coefficient is

$$\hat{\gamma}_h = \Pi_{[0, \gamma_{\text{max}}]}(\tilde{\gamma}_h) , \quad (73)$$

where the Euclidean projection is

$$\Pi_{[0, \gamma_{\text{max}}]}(x) = \min \{ \gamma_{\text{max}}, \max \{ 0, x \} \} . \quad (74)$$

The lower boundary suppresses a residual correction whose calibration association is non-positive, whereas the upper boundary prevents excessive amplification. Only one scalar is estimated per horizon, thereby limiting calibration variance without requiring nonlinear gate optimization. The FastSketch forecast is

$$\hat{y}_{t+h|t}^{\text{FS}} = \hat{y}_{t+h|t}^{\text{base}} + \hat{\gamma}_h \hat{r}_{t+h|t} , \quad h = 1, \dots, H . \quad (75)$$

The reported experiments use $H = 1$. The multi-horizon notation is retained to show that the architecture extends directly by assigning an independent readout column and shrinkage coefficient to each horizon.

2.5. Adaptive Conformal Uncertainty Estimation

The same uncertainty-estimation layer is applied after either S3-Forecaster or S3-FastSketch. Let $\hat{y}_{t|t-1}$ denote the final point forecast. The nonconformity score is

$$s_t = |y_t - \hat{y}_{t|t-1}| . \quad (76)$$

For $H > 1$, a simultaneous score can be defined as

$$s_t^{(H)} = \max_{h=1, \dots, H} |y_{t+h} - \hat{y}_{t+h|t}| . \quad (77)$$

2.5.1. Dynamic Score Buffer

Let W denote the maximum calibration-buffer length. Immediately before forecasting y_t , the score buffer is

$$\mathcal{S}_t = \{s_j : \max(t - W, t_{\text{cal},1}) \leq j \leq t - 1\}, \quad (78)$$

where $t_{\text{cal},1}$ is the first calibration index. Thus,

$$|\mathcal{S}_t| \leq W. \quad (79)$$

Given an effective miscoverage level δ_t , define

$$k_t = \min \{|\mathcal{S}_t|, \lceil (|\mathcal{S}_t| + 1)(1 - \delta_t) \rceil\}. \quad (80)$$

The finite-buffer empirical quantile is

$$q_t(\delta_t) = s_{(k_t)}(\mathcal{S}_t), \quad (81)$$

where $s_{(k_t)}(\mathcal{S}_t)$ is the k_t -th order statistic. The adaptive conformal interval is

$$\widehat{C}_t = [\widehat{y}_{t|t-1} - q_t(\delta_t), \widehat{y}_{t|t-1} + q_t(\delta_t)]. \quad (82)$$

2.5.2. Sequential Miscoverage Update

After observing y_t , define

$$M_t = \mathbb{I}\{y_t \notin \widehat{C}_t\}. \quad (83)$$

Let δ^* be the target miscoverage, with $\delta^* = 0.10$ for nominal 90% coverage. The effective miscoverage is updated as

$$\delta_{t+1} = \Pi_{[\delta_{\min}, \delta_{\max}]} [\delta_t + \eta_{\text{ACI}} (\delta^* - M_t)], \quad (84)$$

where

$$\Pi_{[\delta_{\min}, \delta_{\max}]}(x) = \min \{\delta_{\max}, \max \{\delta_{\min}, x\}\}. \quad (85)$$

When a miss occurs, $M_t = 1$, the update decreases δ_{t+1} and therefore raises the next empirical quantile, producing a wider interval. When coverage is repeatedly achieved, δ_t can increase gradually, reducing excessive width. Finally,

$$\mathcal{S}_{t+1} = \text{Tail}_W(\mathcal{S}_t \cup \{s_t\}). \quad (86)$$

This construction does not assume a Gaussian residual distribution. Its distribution-free interpretation concerns the absence of a parametric error law under conformal validity conditions. Under serial dependence and temporal drift, ACI targets adaptive long-run miscoverage control; it does not imply exact conditional coverage for every individual forecast or every short test block.

2.6. Model Complexity

For S3-Forecaster, the number of learned readout coefficients is approximately

$$P_{\text{S3}} = H(1 + R + p) + 2, \quad (87)$$

where the final two parameters correspond to the OOB gate. The fixed matrices $\mathbf{W}_{\text{in}}$ and $\mathbf{W}_{\text{res}}$ are excluded because they are not trained.

For S3-FastSketch,

$$P_{\text{FS}} = HD_{\phi} + H, \quad (88)$$

where the final H coefficients are the calibration shrinkage factors. For a dense reservoir, state propagation requires

$$\mathcal{O}(NR^2) \quad (89)$$

operations, which can be reduced using a sparse recurrent matrix. FastSketch feature construction requires approximately

$$\mathcal{O}(ND_{\phi}). \quad (90)$$

Both methods avoid iterative backpropagation through a nonlinear sequence model. Their learned readouts are obtained through convex regularized objectives, while FastSketch additionally avoids nonlinear gate calibration.

3. Experimental Setup

3.1. Datasets and Natural Small-Data Cohorts

The empirical evaluation comprises four public monthly forecasting collections—M3 Monthly, M4 Monthly, Tourism Monthly, and CIF 2016—and an anonymized Industrial Construction-Materials Demand dataset, hereafter denoted ICMD. The same experimental procedure is applied independently to every dataset.

The public collections cover heterogeneous forecasting domains, including macroeconomic, demographic, tourism-demand, banking, and general competition series. ICMD contains 61,395 transactional records, 47 original variables, 555 product references, and 62 monthly periods between January 2020 and February 2025. For each eligible product reference i , the monthly target is constructed as

$$y_{i,t} = \sum_{j \in \mathcal{T}_{i,t}} \text{SalesValue}_j, \quad (91)$$

where $\mathcal{T}_{i,t}$ denotes the set of valid positive-value transactions associated with product i during month t . In contrast to a single-reference case study, ICMD is evaluated as a multi-series cohort containing all product-level histories that satisfy the predefined eligibility criteria.

The primary benchmark is restricted to naturally short series. A series is eligible when it has monthly frequency, finite observations, a sufficient unmodified history for chronological training and calibration, and fewer than 200 available pre-test observations. Longer histories are excluded rather than artificially truncated. Consequently, the main benchmark evaluates naturally occurring small-data conditions. Controlled reductions of the available history are introduced only in the Data Efficiency experiment described in Section 3.9.

When a collection contains more than 200 eligible series, at most 200 are selected using a deterministic hash-based ordering with seed 42. This selection is performed before model execution and is independent of forecasting performance. All eligible series are retained when the collection contains fewer than 200 candidates. Series identifiers, original lengths, selected development series, and exclusion reasons are stored with the experimental outputs.

3.2. Development, Failure, and Ranking Cohorts

Let $\mathcal{E}_d$ denote the eligible series in dataset d , and let $\mathcal{D}_d \subset \mathcal{E}_d$ denote the series reserved for dataset-level model selection. Development series are excluded from the final benchmark.

Model failures are not silently removed from the experimental record. For each model m , let

$$\mathcal{S}_{m,d} = \{i \in \mathcal{E}_d \setminus \mathcal{D}_d : m \text{ completes the evaluation of series } i\} \quad (92)$$

denote its successful evaluation cohort. The common cohort used for the primary dataset-level ranking is

$$\mathcal{C}_d = (\mathcal{E}_d \setminus \mathcal{D}_d) \cap \bigcap_{m \in \mathcal{M}} \mathcal{S}_{m,d}, \quad (93)$$

where $\mathcal{M}$ is the complete model set. Thus, every model is ranked using exactly the same series. Model-specific failures, exception messages, and failure rates are reported separately, together with the sizes of $\mathcal{E}_d$, $\mathcal{D}_d$, $\mathcal{S}_{m,d}$, and $\mathcal{C}_d$. Dataset-level metrics are not pooled across collections with different scales.

3.3. Reference Models

The proposed S3-Forecaster and S3-FastSketch are compared with reference models spanning complementary forecasting paradigms:

1. **Statistical forecasters:** Naive, Seasonal Naive, AutoRegressive (AR), ARIMA, ETS, and Theta;
2. **Kernel methods:** Kernel Ridge Regression (KRR) and Gaussian Process Regression (GPR);
3. **Lightweight neural forecasters:** EasyTSF-compatible DLinear and NLinear models, a causal one-dimensional CNN, and a unidirectional LSTM;
4. **Time-series foundation models:** Chronos-Bolt-Small, using the `amazon/chronos-bolt-small` checkpoint, and TimesFM-2.5-200M, using `google/timesfm-2.5-200m-pytorch`;

5. Proposed models: S3-Forecaster and S3-FastSketch.

DLinear, NLinear, CNN, and LSTM use chronological input windows and a temporal validation suffix contained entirely within the training data. Neural samples are not shuffled, and early stopping is determined without access to calibration or test observations. Chronos and TimesFM remain frozen and are evaluated in zero-shot mode; no target-domain gradient update is applied to either foundation model.

Prophet is excluded a priori because repeated fitting over several hundred series and rolling evaluation steps produces a disproportionate computational cost. Its exclusion is based on operational feasibility rather than observed forecasting accuracy. ETS and Theta preserve the statistical trend and seasonality reference class that Prophet would otherwise partially duplicate.

3.4. Offline Model Selection and Frozen Reuse

Hyperparameter selection and final evaluation are treated as disjoint computational phases. For model m and dataset d , let $\mathcal{G}_{m,d}$ denote the predefined dataset-level search space. The selected configuration is

$$\boldsymbol{\theta}_{m,d}^* = \arg \min_{\boldsymbol{\theta} \in \mathcal{G}_{m,d}} \frac{1}{|\mathcal{D}_d|} \sum_{i \in \mathcal{D}_d} \mathcal{L}_{\text{val}}(m, \boldsymbol{\theta}; \mathbf{y}^{(i)}). \quad (94)$$

For the S3 models, point-forecast selection uses three chronological validation folds of six observations, 100 optimization trials, and epsilon-stabilized MAPE as the objective. Uncertainty hyperparameters are selected in a subsequent 60-trial study by minimizing

$$\mathcal{J}_{\text{UQ}} = \text{MSIS} + \lambda_{\text{cov}} [\max\{0, 0.90 - \text{ECP}\}]^2. \quad (95)$$

After selection, $\boldsymbol{\theta}_{m,d}^*$ is serialized and reused without modification in the main benchmark, ablation, Shock/Non-Shock, Data Efficiency, and Multi-prior experiments. No per-series or test-dependent retuning is permitted. The one-time HPO cost is reported separately and is not included in operational forecasting time.

3.5. Chronological Splitting and Rolling Evaluation

For a series $\mathbf{y}^{(i)} = (y_1^{(i)}, \dots, y_{N_i}^{(i)})$, the chronological boundaries are

$$n_{\text{tr}}^{(i)} = \lfloor 0.64N_i \rfloor, \quad (96)$$

$$n_{\text{dev}}^{(i)} = \lfloor 0.80N_i \rfloor, \quad (97)$$

$$n_{\text{cal}}^{(i)} = n_{\text{dev}}^{(i)} - n_{\text{tr}}^{(i)}, \quad (98)$$

$$n_{\text{te}}^{(i)} = N_i - n_{\text{dev}}^{(i)}. \quad (99)$$

These boundaries define disjoint training, calibration, and test blocks:

$$\mathcal{I}_{\text{tr}}^{(i)} = \{1, \dots, n_{\text{tr}}^{(i)}\}, \quad (100)$$

$$\mathcal{I}_{\text{cal}}^{(i)} = \{n_{\text{tr}}^{(i)} + 1, \dots, n_{\text{dev}}^{(i)}\}, \quad (101)$$

$$\mathcal{I}_{\text{te}}^{(i)} = \{n_{\text{dev}}^{(i)} + 1, \dots, N_i\}. \quad (102)$$

All models are evaluated under one-step-ahead rolling forecasting with horizon $H = 1$ and rolling step equal to one. For every $t \in \mathcal{I}_{\text{te}}^{(i)}$, the forecast and interval satisfy

$$\hat{y}_{t|t-1}, \hat{L}_{t|t-1}, \hat{U}_{t|t-1} \in \mathcal{F}_{t-1}, \quad (103)$$

where $\mathcal{F}_{t-1}$ contains only information available before observing y_t . Once the prediction has been recorded, y_t may be used to update the causal state and conformal buffer for step $t + 1$. Dataset-level hyperparameters remain frozen throughout the test block.

Preprocessing statistics are estimated exclusively from the training block. Forecasts and interval bounds are inverse-transformed before evaluation, and all reported errors are computed on the original data scale.

3.6. Common Adaptive Conformal Wrapper

To ensure a homogeneous uncertainty comparison, the same Adaptive Conformal Inference wrapper is applied to every point forecaster, including statistical, kernel, neural, foundation, and S3 models. Native intervals from probabilistic baselines are not used in the primary comparison; they may be reported separately as a supplementary sensitivity analysis.

Let

$$s_t = |y_t - \hat{y}_{t|t-1}| \quad (104)$$

denote the absolute nonconformity score, and let q_t be the empirical $1 - \alpha_t$ quantile of the causally available conformity buffer. The prediction interval is

$$\hat{C}_t = [\hat{y}_{t|t-1} - q_t, \hat{y}_{t|t-1} + q_t]. \quad (105)$$

For target miscoverage $\alpha = 0.10$, the adaptive level is updated as

$$\alpha_{t+1} = \Pi_{\mathcal{A}} \left[\alpha_t + \gamma \left(\alpha - \mathbb{I}\{y_t \notin \hat{C}_t\} \right) \right], \quad (106)$$

where γ is the frozen adaptation step and $\Pi_{\mathcal{A}}$ projects the update onto the admissible range. This protocol targets a nominal coverage of 90% without allowing future test observations to affect the current interval.

3.7. Component-Wise Ablation Study

Ablation experiments reuse the frozen point and uncertainty configurations. Each ablated variant is evaluated on the same chronological partitions and paired series as its corresponding complete model.

For S3-Forecaster, the evaluated variants remove the residual adapter, reservoir states, autoregressive lags, volatility gate, or adapter selector. A static conformal variant is obtained by setting the ACI adaptation step to zero.

For S3-FastSketch, the variants remove the residual adapter, autoregressive lags, EMA residual features, volatility features, momentum features, causal convolutions, seasonal features, calendar features, the complete multiscale sketch, calibration shrinkage, or the adapter selector. A static conformal variant is also included.

For series i , model family a , and ablation v , the paired loss difference is

$$\Delta\mathcal{L}_{a,v}^{(i)} = \mathcal{L}_{a,v}^{(i)} - \mathcal{L}_{a,\text{full}}^{(i)}. \quad (107)$$

Positive values indicate degradation relative to the complete architecture. Ablation summaries are calculated over the common set of series completed by the full model and all compared variants.

3.8. Shock and Non-Shock Analysis

Series are assigned to Shock or Non-Shock regimes using training observations only. For monthly seasonality $m = 12$, the seasonal increment is

$$d_t^{(i)} = y_t^{(i)} - y_{t-m}^{(i)}, \quad t \in \mathcal{I}_{\text{tr}}^{(i)}. \quad (108)$$

A robust location and scale are estimated as

$$c_i = \text{median}_t \left(d_t^{(i)} \right), \quad (109)$$

$$\sigma_i = 1.4826 \text{ median}_t \left| d_t^{(i)} - c_i \right|. \quad (110)$$

When the MAD scale is numerically degenerate, IQR/1.349 is used as a fallback. The robust shock score is

$$z_t^{(i)} = \frac{\left| d_t^{(i)} - c_i \right|}{\max(\sigma_i, \varepsilon_\sigma)}. \quad (111)$$

Observations satisfying $z_t^{(i)} > 3.5$ are flagged, and consecutive flags are collapsed into a single event. A series is assigned to the Shock regime when at least one such event is present in its training block; otherwise, it is assigned to the Non-Shock regime. The label is fixed before test evaluation.

This experiment is therefore a series-stratified robustness analysis rather than a retrospective classification of test observations. For each model, the regime penalty is summarized through

$$R_m^{\text{shock}} = \frac{\tilde{\mathcal{L}}_{m,\text{shock}}}{\tilde{\mathcal{L}}_{m,\text{nonshock}} + \varepsilon}, \quad (112)$$

where $\tilde{\mathcal{L}}$ denotes the median per-series loss within the corresponding regime. Point accuracy, empirical coverage, interval width, and MSIS are reported separately for Shock and Non-Shock series.

3.9. Data Efficiency Analysis

Data efficiency is evaluated at retained-history fractions

$$\mathcal{R} = \{0.25, 0.50, 0.75, 1.00\}. \quad (113)$$

For each fraction r , only the most recent $\lceil rN_i^{\text{pre}} \rceil$ observations of the pre-test history are retained, while the test block, random seed, model configuration, and series identifier remain unchanged. The $r = 1$ condition reuses the full-history benchmark.

A common admissible cohort is determined before executing any fraction. A series is retained only when the smallest history fraction satisfies the seasonal and context-length requirements of every compared model:

$$\lceil r_{\min} N_i^{\text{pre}} \rceil \geq \max \{2m, c_{\max} + H + 1\}, \quad (114)$$

where $c_{\max}$ is the maximum required input context. All reference and proposed models are evaluated in this analysis. Failed model-fraction cells are explicitly reported and are excluded from comparative curves rather than being silently aggregated.

3.10. Multi-Prior Robustness

The dependence of the S3 architectures on the baseline prior is assessed without re-optimizing their residual adapters. Let p^* denote the prior selected during HPO. The evaluated prior set is

$$\mathcal{P} = \{p^*, \text{causal rolling mean, Seasonal Naive, ETS, Theta, Chronos, TimesFM}\}. \quad (115)$$

For every $p \in \mathcal{P}$, only the prior definition and its required fixed parameters are replaced. Reservoir size, residual lags, sketch features, regularization, shrinkage, gate configuration, calibration split, and ACI hyperparameters remain frozen. Prophet is not included because its repeated rolling execution would impose a disproportionately high computational burden.

The sensitivity of architecture a to prior p is measured by

$$\Delta \mathcal{L}_{a,p}^{(i)} = \mathcal{L}_{a,p}^{(i)} - \mathcal{L}_{a,p^*}^{(i)}. \quad (116)$$

The analysis reports the median and interquartile range of these differences, the percentage of improved series, the worst-performing prior, between-prior dispersion, and the number of failed evaluations.

3.11. Evaluation Metrics and Dataset-Level Aggregation

Let $\mathcal{C}_d$ denote the common evaluation cohort for dataset d , and let $\mathcal{T}_i$ be the test set of series i , with $H_i = |\mathcal{T}_i|$. For each $t \in \mathcal{T}_i$, $y_{i,t}$ denotes the observed value, $\hat{y}_{i,t}$ the one-step-ahead point forecast, and $[\hat{L}_{i,t}, \hat{U}_{i,t}]$ the corresponding prediction interval. The forecast error is defined as

$$e_{i,t} = y_{i,t} - \hat{y}_{i,t}. \quad (117)$$

All metrics are first computed independently for each model–series pair. Dataset-level statistics are subsequently obtained across the common cohort $\mathcal{C}_d$, thereby assigning equal weight to every series rather than pooling observations from series with different test lengths.

3.11.1. Point-Forecast Metrics

The epsilon-stabilized Mean Absolute Percentage Error is

$$\text{MAPE}_\varepsilon^{(i)} = \frac{100}{H_i} \sum_{t \in \mathcal{T}_i} \frac{|e_{i,t}|}{\max(|y_{i,t}|, \varepsilon)}, \quad \varepsilon = 10^{-5}. \quad (118)$$

The symmetric Mean Absolute Percentage Error is defined as

$$\text{sMAPE}_\varepsilon^{(i)} = \frac{100}{H_i} \sum_{t \in \mathcal{T}_i} \frac{2|e_{i,t}|}{\max(|y_{i,t}| + |\hat{y}_{i,t}|, \varepsilon)}. \quad (119)$$

To construct a scale-independent absolute error, let $\mathbf{y}_i^{\text{pre}} = (y_{i,1}, \dots, y_{i,n_i^{\text{pre}}})$ denote the complete pre-test history available for fitting and calibration. The seasonal naive scaling factor is

$$s_i = \frac{1}{n_i^{\text{pre}} - m} \sum_{t=m+1}^{n_i^{\text{pre}}} |y_{i,t} - y_{i,t-m}|, \quad (120)$$

where $m = 12$ for the monthly datasets. The Mean Absolute Scaled Error is then

$$\text{MASE}^{(i)} = \frac{\frac{1}{H_i} \sum_{t \in \mathcal{T}_i} |e_{i,t}|}{s_i}. \quad (121)$$

A value $\text{MASE} < 1$ indicates that the evaluated model produces a lower mean absolute error than the seasonal naive reference associated with the scaling denominator.

The Root Mean Squared Error is computed on the original data scale as

$$\text{RMSE}^{(i)} = \sqrt{\frac{1}{H_i} \sum_{t \in \mathcal{T}_i} e_{i,t}^2}. \quad (122)$$

MAPE and sMAPE are reported as percentages, MASE provides a scale-free comparison against the seasonal naive reference, and RMSE retains sensitivity to large forecast errors. Consequently, these measures are interpreted jointly rather than as interchangeable estimates of predictive performance.

3.11.2. Prediction-Interval Metrics

All uncertainty metrics are evaluated at nominal coverage $1 - \alpha = 0.90$, with $\alpha = 0.10$. The Empirical Coverage Probability for series i is

$$\text{ECP}^{(i)} = \frac{1}{H_i} \sum_{t \in \mathcal{T}_i} \mathbb{I} \left\{ \hat{L}_{i,t} \leq y_{i,t} \leq \hat{U}_{i,t} \right\}. \quad (123)$$

The mean prediction-interval width is

$$\text{MW}^{(i)} = \frac{1}{H_i} \sum_{t \in \mathcal{T}_i} \left(\hat{U}_{i,t} - \hat{L}_{i,t} \right). \quad (124)$$

Because coverage alone can be increased by constructing arbitrarily wide intervals, interval calibration and sharpness are evaluated jointly through the interval score. For observation t , the interval score is

$$\begin{aligned} \text{IS}_{\alpha,i,t} &= \left(\hat{U}_{i,t} - \hat{L}_{i,t} \right) \\ &\quad + \frac{2}{\alpha} \left(\hat{L}_{i,t} - y_{i,t} \right) \mathbb{I} \left\{ y_{i,t} < \hat{L}_{i,t} \right\} \\ &\quad + \frac{2}{\alpha} \left(y_{i,t} - \hat{U}_{i,t} \right) \mathbb{I} \left\{ y_{i,t} > \hat{U}_{i,t} \right\}. \end{aligned} \quad (125)$$

The Mean Scaled Interval Score is obtained by normalizing the mean interval score using the same seasonal naive scale employed by MASE:

$$\text{MSIS}^{(i)} = \frac{\frac{1}{H_i} \sum_{t \in \mathcal{T}_i} \text{IS}_{\alpha,i,t}}{s_i}. \quad (126)$$

Lower MSIS values indicate sharper intervals with smaller penalties for observations falling outside the interval. ECP is therefore interpreted relative to the nominal value of 0.90 and jointly with mean width and MSIS; overcoverage is not considered beneficial when it is obtained through excessive interval inflation.

3.11.3. Dataset-Level Aggregation

For any per-series metric $x_{j,i}$ associated with model j and series i , the dataset-level median is

$$\tilde{x}_{j,d} = Q_{0.50} \left(\{x_{j,i} : i \in \mathcal{C}_d\} \right), \quad (127)$$

and its interquartile range is

$$\text{IQR}(x_{j,d}) = Q_{0.75} \left(\{x_{j,i} : i \in \mathcal{C}_d\} \right) - Q_{0.25} \left(\{x_{j,i} : i \in \mathcal{C}_d\} \right). \quad (128)$$

MAPE, sMAPE, MASE, RMSE, mean interval width, MSIS, trainable parameters, and elapsed time are summarized using the median and IQR. Since coverage is a bounded probability with a direct nominal target, dataset-level ECP is calculated as the unweighted mean of the per-series coverage values:

$$\overline{\text{ECP}}_{j,d} = \frac{1}{|\mathcal{C}_d|} \sum_{i \in \mathcal{C}_d} \text{ECP}_{j,i}. \quad (129)$$

The reported number of series is

$$n_{\text{series},d} = |\mathcal{C}_d|, \quad (130)$$

which is identical across models in the primary ranking because all comparisons use the common successful cohort.

Let $P_{j,i}$ denote the number of target-domain trainable parameters and let $\tau_{j,i}$ denote the elapsed time required for final fitting and rolling prediction. Computational summaries are defined as

$$\tilde{P}_{j,d} = \text{median}_{i \in \mathcal{C}_d} P_{j,i}, \quad (131)$$

$$\tilde{\tau}_{j,d} = \text{median}_{i \in \mathcal{C}_d} \tau_{j,i}. \quad (132)$$

Frozen foundation-model weights and fixed random projections are excluded from $P_{j,i}$ because they are not adjusted using the target series. Their total parameter footprint is reported separately. Elapsed time excludes environment installation, checkpoint downloading, and offline hyperparameter optimization.

Finally, the dataset-specific ordinal ranking is obtained by sorting the models in ascending order of median MAPE_ε :

$$\text{Rank}_{j,d} = \text{rank}_\uparrow \left(\widetilde{\text{MAPE}_{\varepsilon,j,d}} \right). \quad (133)$$

The ranking is interpreted together with sMAPE, MASE, RMSE, ECP, mean width, and MSIS, since no individual metric fully characterizes point accuracy, robustness to scale, and uncertainty quality.

Operational runtime includes final fitting and rolling forecasting after loading the frozen configuration. HPO, checkpoint downloading, and environment installation are excluded and reported separately. Trainable-parameter counts include coefficients estimated from the target series but exclude frozen foundation-model weights and fixed random feature matrices. The total parameter footprint of frozen pretrained models is reported separately to avoid interpreting zero target-domain trainable parameters as zero computational complexity.

3.12. Reproducibility Controls

All experiments use fixed random seeds and repository commits. Dataset identifiers, natural series lengths, development exclusions, common cohorts, frozen configurations, model failures, per-series forecasts, and per-series metrics are persisted before aggregation. Neural models are trained without temporal shuffling, and no test observation is used for preprocessing, model selection, early stopping, gate calibration, or conformal initialization. These controls ensure that all reported rankings correspond to a common causal evaluation protocol rather than to model-specific subsets or retrospectively selected histories.

Table 1: Phase-1 benchmark summary based on median MAPE. Only completed model–dataset evaluations are ranked.

Dataset	S3-FastSketch		S3-Forecaster		Overall best		Best non-S3	
	Rank	MAPE (%)	Rank	MAPE (%)	Model	MAPE (%)	Model	MAPE (%)
CIF	1	5.00	2	5.13	S3-FastSketch	5.00	ETS	5.19
M3	11	8.80	6	7.04	ETS	5.44	ETS	5.44
M4	7	3.97	8	4.05	GPR	2.97	GPR	2.97
Tourism	1	16.64	4	18.09	S3-FastSketch	16.64	ETS	17.32

Note: Ranks are conditional on the available phase-1 results and must not be interpreted as final rankings until the missing model–dataset runs are completed.

4. Results and Discussion

4.1. Phase-1 Evidence Scope and Reporting Strategy

This section reports the first completed experimental block and is intentionally restricted to the evidence contained in the current exports. The benchmark file contains 52 unique model–dataset evaluations over CIF, M3, M4, and Tourism. All 16 candidate models are available for CIF and M3, whereas M4 and Tourism currently contain 10 models and do not yet include CNN, LSTM, DLinear, NLinear, Chronos-Bolt-Small, or TimesFM-2.5-200M. Consequently, the ranks reported for M4 and Tourism are conditional on the completed runs and are not final cross-model rankings.

The error distributions are strongly right-skewed. In the data-efficiency export, the ratio between mean and median MAPE reaches several orders of magnitude for some cells, consistent with the instability of percentage errors when individual targets approach zero. Median MAPE is therefore used as the primary descriptive statistic, with the interquartile range (IQR), MASE, and empirical coverage probability (ECP) used as complementary evidence. This choice prevents a small number of pathological percentage errors from dominating the cross-series interpretation, but it does not remove the known limitations of MAPE.

The anonymized Industrial Construction-Materials Demand (ICMD) application remains the primary engineering use case. Its source contains 61,395 transactional records, 47 variables, 555 product references, and 62 monthly periods between January 2020 and February 2025. The evaluated product series contains 48 monthly observations, with the final 10 months reserved for test. The current phase-1 exports do not contain the finalized ICMD benchmark rows; therefore, no ICMD accuracy or shock-robustness claim is made in this draft. Those results will be inserted without altering the analyses reported below.

4.2. Overall Benchmark Performance

Table 1 and Fig. 1 show that the relative performance of the S3 family is dataset dependent rather than uniformly dominant. This result is important because the proposed architectures are intended as controlled residual adapters, not as universal replacements for every statistical forecaster.

On CIF, S3-FastSketch ranks first with a median MAPE of 4.996%, followed by S3-Forecaster at 5.133%. The strongest completed non-S3 model is ETS at 5.189%, so the improvement of S3-FastSketch over that

reference is 3.7% in relative terms. The two S3 variants also precede the zero-shot foundation models and the four trained neural baselines on this dataset. In particular, DLinear, NLinear, CNN, and LSTM obtain median MAPEs of 10.166%, 9.001%, 11.155%, and 16.993%, respectively. This ordering is consistent with the small-data motivation: increasing end-to-end trainable capacity does not compensate automatically for limited series length.

Tourism provides a second favorable case. S3-FastSketch ranks first among the 10 completed models with a median MAPE of 16.643%, compared with 17.319% for ETS and 17.964% for Theta. This corresponds to a 3.9% relative reduction with respect to the best completed non-S3 comparator. S3-Forecaster is fourth at 18.091%. The stronger result of the deterministic sketch variant suggests that a fixed multiscale residual representation can be preferable to a richer reservoir when the available signal is heterogeneous across many short series. This interpretation remains provisional because the six neural and foundation runs for Tourism are not yet present.

The results on M3 and M4 delimit the claim. On M3, ETS and Theta lead with median MAPEs of 5.442% and 5.489%. S3-Forecaster ranks sixth at 7.038%, while S3-FastSketch ranks eleventh at 8.797%. Nevertheless, S3-Forecaster remains more accurate than Chronos-Bolt-Small, TimesFM-2.5-200M, DLinear, NLinear, CNN, and LSTM in the completed M3 comparison. On M4, the two S3 variants rank seventh and eighth with median MAPEs of 3.968% and 4.046%, whereas GPR leads at 2.972%. Thus, the present evidence supports complementarity: the S3 family is most useful when the frozen prior leaves a learnable residual structure, while parsimonious statistical or kernel models remain stronger when that condition is not satisfied.

Point accuracy and interval calibration do not move uniformly. S3-FastSketch and S3-Forecaster obtain mean ECP values of 0.909 and 0.955 on CIF, but the corresponding values are 0.837 and 0.933 on M3, 0.949 and 0.970 on M4, and 0.848 and 0.887 on Tourism. Relative to the nominal 0.90 target, the current intervals therefore range from under-coverage to substantial over-coverage. The adaptive conformal layer should consequently be presented as an online calibration mechanism rather than as evidence of uniform finite-sample calibration across all datasets. Shock-stratified interval diagnostics and the pending ICMD run are required before making a stronger uncertainty claim.

4.3. Data Efficiency Under Reduced Histories

Figure 2 and Table 2 summarize the available reduced-history runs. The trajectories are non-monotonic. For S3-FastSketch on CIF, median MAPE falls from 5.983% at 25% history to 4.894% at 75%, and remains 4.996% with the full history. S3-Forecaster is stable between 50% and 100% history, with median MAPE between 4.972% and 5.133%. On M3, S3-Forecaster reaches 6.110% at 25% and 5.178% at 50%, but increases to approximately 7.0% at 75% and 100%. A similar non-monotonic pattern occurs on M4: the 25% values are 3.441% for S3-FastSketch and 2.865% for S3-Forecaster, while the corresponding full-history values are 3.968% and 4.046%.

Table 2: Median MAPE (%) as a function of the available training history.

Dataset	Model	25% (<i>n</i>)	50% (<i>n</i>)	75% (<i>n</i>)	100% (<i>n</i>)
CIF	S3-FastSketch	5.98 (47)	5.55 (61)	4.89 (65)	5.00 (65)
CIF	S3-Forecaster	–	4.97 (53)	5.13 (61)	5.13 (65)
M3	S3-FastSketch	7.64 (149)	9.68 (199)	9.14 (199)	8.80 (199)
M3	S3-Forecaster	6.11 (143)	5.18 (155)	7.02 (199)	7.04 (199)
M4	S3-FastSketch	3.44 (153)	5.29 (199)	4.36 (199)	3.97 (199)
M4	S3-Forecaster	2.87 (145)	4.31 (199)	3.94 (199)	4.05 (199)
Tourism	S3-FastSketch	–	–	–	16.64 (199)
Tourism	S3-Forecaster	–	–	–	18.09 (199)

Note: Each cell reports median MAPE followed by the number of successfully evaluated series in parentheses. Cohort size changes across fractions and, in some cases, across models; the trajectories are therefore descriptive rather than fixed-cohort learning curves. “–” denotes an unavailable run.

These values show that both architectures can remain operational with sharply reduced histories, but they do not yet establish a monotonic learning curve. The number of successfully evaluated series changes across fractions: for example, the CIF S3-FastSketch cohort increases from 47 series at 25% history to 65 at full history, while the M4 S3-Forecaster cohort increases from 145 to 199. The current curves therefore combine a history-length effect with a cohort composition effect. They should be interpreted as phase-1 feasibility evidence, not as a causal estimate of the value of additional observations. The final analysis should rerun every fraction on the intersection of series that is valid for all models and fractions, and should report paired uncertainty for the resulting changes.

The non-monotonic behavior is nevertheless scientifically informative. More history is not automatically more useful under non-stationarity: older regimes may dilute the local residual structure relevant to the test block, while shorter histories may reduce model variance. This mechanism is compatible with the design of S3, but it cannot be distinguished from cohort selection using the aggregated export alone. No such causal attribution is made here.

4.4. Shock Versus Non-Shock Behaviour

All 52 completed model–dataset combinations contain both shock and non-shock summaries. Figure 3 reports the ratio between shock and non-shock median MAPE, while Table 3 gives the absolute values for the S3 family. A ratio above one indicates that accuracy deteriorates on the observations classified as shocks.

The regime split exposes a consistent challenge: most models deteriorate under shocks, but the magnitude depends strongly on the dataset. On CIF, shock MAPE is 7.771% for S3-FastSketch and 7.940% for S3-Forecaster, approximately twice their non-shock errors of 3.821% and 4.045%. These absolute shock errors are close to ETS at 7.500%, while remaining below the shock errors of the neural and foundation baselines. On Tourism, both S3 variants show a smaller relative penalty: their shock/non-shock ratios are 1.25 and 1.26, with shock MAPEs of 18.666% and 18.721%. ETS is slightly better in absolute terms at 18.367%,

Table 3: Shock versus non-shock median MAPE for the S3 family.

Dataset	Model	Non-shock (% , n)	Shock (% , n)	Ratio
CIF	S3-FastSketch	3.82 (40)	7.77 (25)	2.03
CIF	S3-Forecaster	4.05 (40)	7.94 (25)	1.96
M3	S3-FastSketch	4.85 (88)	13.08 (111)	2.70
M3	S3-Forecaster	2.95 (88)	9.85 (111)	3.34
M4	S3-FastSketch	1.73 (83)	6.59 (116)	3.81
M4	S3-Forecaster	1.82 (83)	6.46 (116)	3.54
Tourism	S3-FastSketch	14.92 (31)	18.67 (168)	1.25
Tourism	S3-Forecaster	14.90 (31)	18.72 (168)	1.26

Note: The ratio is shock MAPE divided by non-shock MAPE. Values above one indicate degradation during shock observations.

whereas Theta reaches 19.446%.

M3 and M4 are more difficult. On M3, S3-FastSketch rises from 4.849% to 13.080% and S3-Forecaster from 2.951% to 9.848%. Their ratios of 2.70 and 3.34 are lower than those of AR, ARIMA, ETS, and Naive, but S3 does not achieve the lowest absolute shock error; Theta reaches 8.158% and ETS 8.740%. On M4, the S3 shock errors are 6.589% and 6.457%, compared with 4.858% for ETS and 5.002% for AR. The shock-aware components therefore attenuate neither every shock nor every dataset equally.

The two Naive ratios below one on CIF and Tourism should not be interpreted as successful shock adaptation. In both cases the non-shock baseline error is already high, so the ratio is reduced by a weak denominator reference. This is why the absolute-regime table and the relative-penalty figure are reported together. The current aggregate CSV also does not contain per-series paired losses for statistical testing. The final manuscript should add confidence intervals or paired resampling before claiming significant shock robustness.

4.5. Ablation Evidence and Mechanistic Interpretation

The ablation export contains 80 unique dataset–model–variant summaries and uses the same number of evaluated series within each dataset–model block. Figure 4 reports the relative change in aggregated median MAPE after removing one component. Positive values favor the full model; negative values indicate that the ablated variant has lower median error. Table 4 lists the largest favorable and adverse changes for each block.

Two mechanisms receive direct support. First, removing shrinkage from S3-FastSketch increases median MAPE by 24.5% on CIF and 16.6% on Tourism. This is consistent with shrinkage acting as a variance-control mechanism rather than as a cosmetic regularizer. Second, replacing the adapted forecast with the prior-only configuration is especially damaging on M4: median MAPE increases by 44.1% for S3-FastSketch and 41.3% for S3-Forecaster. Prior-only also increases median MAPE by 12.1% and 10.5% on M3. These results show that residual adaptation can add substantial value when the residual signal is exploitable.

Table 4: Largest favorable and adverse paired median MAPE changes in the phase-1 ablation export.

Dataset	Model	Most favorable removal	Change (%)	Most adverse removal	Change (%)
CIF	S3-FastSketch	No EMA residual	-0.6	No shrinkage	24.5
CIF	S3-Forecaster	Prior only	-3.3	No AR lags	7.6
M3	S3-FastSketch	No multiscale sketch	-15.3	Prior only	12.1
M3	S3-Forecaster	No AR lags	-1.2	Prior only	10.5
M4	S3-FastSketch	No causal convolution	-2.5	Prior only	44.1
M4	S3-Forecaster	No volatility gate	-1.7	Prior only	41.3
Tourism	S3-FastSketch	No volatility features	-0.6	No shrinkage	16.6
Tourism	S3-Forecaster	No AR lags	-6.1	No volatility gate	1.8

Note: Changes compare each variant’s aggregated median MAPE with the corresponding full-model median. Negative values favor the ablated variant; positive values favor the full model. This descriptive summary does not establish statistical significance.

The ablations also reveal components that are not universally beneficial. Removing the multiscale sketch reduces the aggregated median MAPE of S3-FastSketch by 15.3% on M3, while removing autoregressive lags reduces the S3-Forecaster median by 6.1% on Tourism. Smaller favorable removals occur in several other blocks. This is not evidence against the overall architecture; it shows that a fixed feature inventory can inject variance when the corresponding residual pattern is absent. The adapter selector and shrinkage mechanisms are therefore central to the S3 argument: the architecture should be judged by its ability to exploit useful residual structure while suppressing unsupported corrections, not by an assumption that every feature group must help on every dataset.

4.6. Implications, Limitations, and Pending Evidence

The phase-1 results support a focused engineering claim. Low-capacity, gradient-free residual adaptation can outperform statistical, trained neural, and frozen foundation references on selected small-data collections, as shown most clearly by the first-place S3-FastSketch results on CIF and Tourism. The M3 and M4 outcomes prevent a stronger universal-superiority claim and instead motivate an explicit model-selection perspective: S3 is most valuable when the prior residual is stable enough to learn, and the prior-only or a classical model remains preferable otherwise.

The neural results currently available for CIF and M3 also clarify the role of LSTM and CNN in the final benchmark. Both models rank below the S3 family on those datasets and require substantially more operational time than the lightweight alternatives on M3. They remain useful as legacy trained-deep controls, but the evidence does not justify making them the center of the main comparison. Retaining them in the completed benchmark or supplementary material is more defensible than removing all trained nonlinear sequence references.

Four limitations define the boundary of this draft. First, M4 and Tourism are missing the six neural and foundation-model runs, so their ranks are conditional. Second, the reduced-history cohorts are not

fixed, preventing a clean causal data-efficiency claim. Third, the shock and ablation exports are aggregated summaries and do not support paired significance tests. Fourth, the final ICMD and multi-prior robustness results are not included in the supplied files. These gaps do not invalidate the present observations, but they must be closed before submission. The final analysis should preserve the same chronological protocol, freeze the evaluated cohort across data fractions, report paired uncertainty, and update this section without extrapolating from missing cells.

Phase-1 benchmark accuracy by dataset




Figure 1: Median MAPE for the completed phase-1 benchmark runs. Blue and orange identify S3-FastSketch and S3-Forecaster, respectively. M4 and Tourism contain only the 10 currently completed models.




Figure 2: Median MAPE of the S3 family as the available training history is reduced. Tourism currently contains only the full-history endpoints. The number of successfully evaluated series changes across some fractions.

Relative accuracy degradation under shock regimes




Figure 3: Ratio of shock to non-shock median MAPE for each completed model. The dashed line denotes equal error in both regimes. Ratios must be read jointly with absolute error because a low ratio can result from poor non-shock performance.




Figure 4: Relative change in aggregated median MAPE for the phase-1 ablations. Positive values indicate degradation after removing the component. Blank cells denote variants that do not apply to an architecture.

5. Conclusions

CRediT authorship contribution statement

Alejandro Patino Bedoya: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Visualization, Project administration, Writing – original draft, Writing – review and editing.

Declaration of competing interest

The author declares that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The M3, M4, Tourism, and CIF 2016 benchmark datasets used in this study are publicly available from the original sources cited in the manuscript. The source code, selected public-series identifiers, truncation lengths, random seeds, serialized hyperparameter configurations, and scripts required to reproduce the public benchmark, ablation, data-efficiency, and multi-prior experiments are openly available at <https://github.com/AlejoPatigno/S3Forecaster-S3FastSketch>.

The Industrial Construction-Materials Demand dataset cannot be made publicly available because the underlying transactional records contain confidential commercial information and are subject to data-use restrictions imposed by the data owner. The manuscript reports the target definition, monthly aggregation procedure, chronological partitions, preprocessing protocol, model configurations, and evaluation procedure required to reproduce the methodological workflow. Access to the proprietary records cannot be guaranteed and would require explicit authorization from the data owner.

Funding

Financial support was provided by Minciencias under the program “Plataforma de analítica de datos como soporte a la gestión de la demanda operacional en la industria manufacturera de la red de aliados de RevloG” code 825223, Universidad Nacional de Colombia

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the author used OpenAI’s ChatGPT to support literature organization, language refinement, and structural editing. After using this tool, the author reviewed, verified, and edited the content as needed and takes full responsibility for the content of the published article.

References

- [1] Y.-J. Park, D. Kim, F. Odermatt, J. Lee, and K.-M. Kim, “A scalable and transferable time series prediction framework for demand forecasting,” arXiv preprint arXiv:2402.19402, 2024. <https://arxiv.org/abs/2402.19402>
- [2] L. Li, Y. Kang, F. Petropoulos, and F. Li, “Feature-based intermittent demand forecast combinations: Bias, accuracy and inventory implications,” arXiv preprint arXiv:2204.08283, 2022. <https://arxiv.org/abs/2204.08283>
- [3] D. Kaur, S. N. Islam, M. A. Mahmud, M. E. Haque, and Z. Y. Dong, “Energy forecasting in smart grid systems: A review of the state-of-the-art techniques,” arXiv preprint arXiv:2011.12598, 2020. <https://arxiv.org/abs/2011.12598>
- [4] T. Zhu, Y. Ran, X. Zhou, and Y. Wen, “A survey of predictive maintenance: Systems, purposes and approaches,” arXiv preprint arXiv:1912.07383, 2019. <https://arxiv.org/abs/1912.07383>
- [5] B. Lim and S. Zohren, “Time-series forecasting with deep learning: A survey,” *Philosophical Transactions of the Royal Society A*, vol. 379, no. 2194, p. 20200209, 2021. <https://doi.org/10.1098/rsta.2020.0209>
- [6] J. M. Rožanec and D. Mladenović, “Reframing demand forecasting: A two-fold approach for lumpy and intermittent demand,” arXiv preprint arXiv:2103.13812, 2021. <https://arxiv.org/abs/2103.13812>
- [7] Y. Gao, Y. Wang, and Q. Wang, “Improving the transferability of time series forecasting with decomposition adaptation,” arXiv preprint arXiv:2307.00066, 2023. <https://arxiv.org/abs/2307.00066>
- [8] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997. <https://doi.org/10.1162/neco.1997.9.8.1735>
- [9] H. Zhou, S. Zhang, J. Peng, S. Zhang, J. Li, H. Xiong, and W. Zhang, “Informer: Beyond efficient transformer for long sequence time-series forecasting,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 12, pp. 11106–11115, 2021. <https://doi.org/10.1609/aaai.v35i12.17325>
- [10] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam, “A time series is worth 64 words: Long-term forecasting with transformers,” in *International Conference on Learning Representations*, 2023. <https://arxiv.org/abs/2211.14730>
- [11] A. F. Ansari, L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. P. Arango, S. Kapoor, J. Zschiegner, D. C. Maddix, H. Wang, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang, “Chronos: Learning the language of time series,” *Transactions on Machine Learning Research*, 2024. <https://arxiv.org/abs/2403.07815>

- [12] A. Das, W. Kong, R. Sen, and Y. Zhou, “A decoder-only foundation model for time-series forecasting,” in *Proceedings of the International Conference on Machine Learning*, 2024. <https://arxiv.org/abs/2310.10688>
- [13] A. Zeng, M. Chen, L. Zhang, and Q. Xu, “Are transformers effective for time series forecasting?,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 9, pp. 11121–11128, 2023. <https://doi.org/10.1609/aaai.v37i9.26317>
- [14] L. Shen, Y. Wei, and Y. Wang, “GBT: Two-stage transformer framework for non-stationary time series forecasting,” arXiv preprint arXiv:2307.08302, 2023. <https://arxiv.org/abs/2307.08302>
- [15] C. Ehrig, B. Sonnleitner, U. Neumann, C. Cleophas, and G. Forestier, “The impact of data set similarity and diversity on transfer learning success in time series forecasting,” arXiv preprint arXiv:2404.06198, 2024. <https://arxiv.org/abs/2404.06198>
- [16] J. D. Hamilton, “Analysis of time series subject to changes in regime,” *Journal of Econometrics*, vol. 45, no. 1–2, pp. 39–70, 1990. [https://doi.org/10.1016/0304-4076\(90\)90093-9](https://doi.org/10.1016/0304-4076(90)90093-9)
- [17] J. D. Hamilton and R. Susmel, “Autoregressive conditional heteroskedasticity and changes in regime,” *Journal of Econometrics*, vol. 64, no. 1–2, pp. 307–333, 1994. [https://doi.org/10.1016/0304-4076\(94\)90067-1](https://doi.org/10.1016/0304-4076(94)90067-1)
- [18] L. Callot and J. T. Kristensen, “Vector autoregressions with parsimoniously time-varying parameters and an application to monetary policy,” arXiv preprint arXiv:1411.0877, 2014. <https://arxiv.org/abs/1411.0877>
- [19] M. M. Fischer, N. Hauzenberger, F. Huber, and M. Pfarrhofer, “General Bayesian time-varying parameter VARs for predicting government bond yields,” arXiv preprint arXiv:2102.13393, 2021. <https://arxiv.org/abs/2102.13393>
- [20] L. Onorante and A. E. Raftery, “Dynamic model averaging in large model spaces using dynamic Occam’s window,” arXiv preprint arXiv:1410.7799, 2014. <https://arxiv.org/abs/1410.7799>
- [21] R. P. Adams and D. J. C. MacKay, “Bayesian online changepoint detection,” arXiv preprint arXiv:0710.3742, 2007. <https://arxiv.org/abs/0710.3742>
- [22] L. J. Wendelberger, J. M. Gray, B. J. Reich, and A. G. Wilson, “Monitoring deforestation using multivariate Bayesian online changepoint detection with outliers,” arXiv preprint arXiv:2112.12899, 2021. <https://arxiv.org/abs/2112.12899>
- [23] J. Lu, A. Liu, F. Dong, F. Gu, J. Gama, and G. Zhang, “Learning under concept drift: A review,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 31, no. 12, pp. 2346–2363, 2020. <https://doi.org/10.1109/TKDE.2018.2876857>

- [24] M. D. L. Tosi and M. Theobald, “OPTWIN: Drift identification with optimal sub-windows,” arXiv preprint arXiv:2305.11942, 2023. <https://arxiv.org/abs/2305.11942>
- [25] Y.-F. Zhang, Q. Wen, X. Wang, W. Chen, L. Sun, Z. Zhang, L. Wang, R. Jin, and T. Tan, “OneNet: Enhancing time series forecasting models under concept drift by online ensembling,” in *Advances in Neural Information Processing Systems*, vol. 36, 2023. <https://arxiv.org/abs/2309.12659>
- [26] Y. Zhang, W. Chen, Z. Zhu, D. Qin, L. Sun, X. Wang, Q. Wen, Z. Zhang, L. Wang, and R. Jin, “Addressing concept shift in online time series forecasting: Detect-then-adapt,” arXiv preprint arXiv:2403.14949, 2024. <https://arxiv.org/abs/2403.14949>
- [27] L. Zhao and Y. Shen, “Proactive model adaptation against concept drift for online time series forecasting,” arXiv preprint arXiv:2412.08435, 2024. <https://arxiv.org/abs/2412.08435>
- [28] A. Casini, “Tests for forecast instability and forecast failure under a continuous record asymptotic framework,” arXiv preprint arXiv:1803.10883, 2018. <https://arxiv.org/abs/1803.10883>
- [29] A. N. Angelopoulos and S. Bates, “Conformal prediction: A gentle introduction,” *Foundations and Trends in Machine Learning*, vol. 16, no. 4, pp. 494–591, 2023. <https://doi.org/10.1561/2200000101>
- [30] T. Bollerslev, “Generalized autoregressive conditional heteroskedasticity,” *Journal of Econometrics*, vol. 31, no. 3, pp. 307–327, 1986. [https://doi.org/10.1016/0304-4076\(86\)90063-1](https://doi.org/10.1016/0304-4076(86)90063-1)
- [31] C. Wang and R. M. Neal, “Gaussian process regression with heteroscedastic or non-Gaussian residuals,” arXiv preprint arXiv:1212.6246, 2012. <https://arxiv.org/abs/1212.6246>
- [32] V. Tolvanen, P. Jylänki, and A. Vehtari, “Approximate inference for nonstationary heteroscedastic Gaussian process regression,” arXiv preprint arXiv:1404.5443, 2014. <https://arxiv.org/abs/1404.5443>
- [33] S. Karlsson, S. Mazur, and H. Nguyen, “Vector autoregression models with skewness and heavy tails,” arXiv preprint arXiv:2105.11182, 2021. <https://arxiv.org/abs/2105.11182>
- [34] G. Shafer and V. Vovk, “A tutorial on conformal prediction,” *Journal of Machine Learning Research*, vol. 9, pp. 371–421, 2008. <https://www.jmlr.org/papers/v9/shafer08a.html>
- [35] M. Zaffran, A. Dieuleveut, O. Féron, Y. Goude, and J. Josse, “Adaptive conformal predictions for time series,” in *Proceedings of the 39th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 162, pp. 25834–25866, 2022. <https://proceedings.mlr.press/v162/zaffran22a.html>
- [36] C. Xu and Y. Xie, “Sequential predictive conformal inference for time series,” in *Proceedings of the 40th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 202, pp. 38707–38727, 2023. <https://arxiv.org/abs/2212.03463>

- [37] I. Gibbs and E. Candès, “Adaptive conformal inference under distribution shift,” in *Advances in Neural Information Processing Systems*, vol. 34, pp. 1660–1672, 2021. <https://arxiv.org/abs/2106.00170>
- [38] I. Gibbs and E. Candès, “Conformal inference for online prediction with arbitrary distribution shifts,” *Journal of Machine Learning Research*, vol. 25, no. 162, pp. 1–36, 2024. <https://www.jmlr.org/papers/v25/22-1218.html>
- [39] V. Cerqueira, L. Torgo, and I. Mozetič, “Evaluating time series forecasting models: An empirical study on performance estimation methods,” *Machine Learning*, vol. 109, pp. 1997–2028, 2020. <https://doi.org/10.1007/s10994-020-05910-7>
- [40] P.-C. Bürkner, J. Gabry, and A. Vehtari, “Approximate leave-future-out cross-validation for Bayesian time series models,” *Journal of Statistical Computation and Simulation*, vol. 90, no. 14, pp. 2499–2523, 2020. <https://doi.org/10.1080/00949655.2020.1783262>
- [41] H. Jaeger and H. Haas, “Harnessing nonlinearity: Predicting chaotic systems and saving energy in wireless communication,” *Science*, vol. 304, no. 5667, pp. 78–80, 2004. <https://doi.org/10.1126/science.1091277>
- [42] L. B. Armenio, E. Terzi, M. Farina, and R. Scattolini, “Echo State Networks: Analysis, training and predictive control,” arXiv preprint arXiv:1902.01618, 2019. <https://arxiv.org/abs/1902.01618>
- [43] M. Lukoševičius and H. Jaeger, “Reservoir computing approaches to recurrent neural network training,” *Computer Science Review*, vol. 3, no. 3, pp. 127–149, 2009. <https://doi.org/10.1016/j.cosrev.2009.03.005>
- [44] W. Maass, T. Natschläger, and H. Markram, “Real-time computing without stable states: A new framework for neural computation based on perturbations,” *Neural Computation*, vol. 14, no. 11, pp. 2531–2560, 2002. <https://doi.org/10.1162/089976602760407955>
- [45] S. Makridakis, E. Spiliotis, and V. Assimakopoulos, “The M4 Competition: 100,000 time series and 61 forecasting methods,” *International Journal of Forecasting*, vol. 36, no. 1, pp. 54–74, 2020. <https://doi.org/10.1016/j.ijforecast.2019.04.014>
- [46] S. Makridakis and M. Hibon, “The M3-Competition: results, conclusions and implications,” *International Journal of Forecasting*, vol. 16, no. 4, pp. 451–476, 2000. [https://doi.org/10.1016/S0169-2070\(00\)00057-1](https://doi.org/10.1016/S0169-2070(00)00057-1)
- [47] G. Athanasopoulos, R. J. Hyndman, H. Song, and D. C. Wu, “The tourism forecasting competition,” *International Journal of Forecasting*, vol. 27, no. 3, pp. 822–844, 2011. <https://doi.org/10.1016/j.ijforecast.2010.04.009>

- [48] M. Štěpnička and M. Burda, “Computational Intelligence in Forecasting Competition: introduction and results,” *International Journal of Forecasting*, vol. 33, no. 1, pp. 1–4, 2017. <https://doi.org/10.1016/j.ijforecast.2016.07.004>